


```
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

Read data

```
data = pd.read_csv('/content/train_v9rqX0R.csv')
data.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614



```
data.shape
```

```
(8523, 12)
```

check for null values

```
data.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          1463
Item_Fat_Content      0
Item_Visibility       0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size          2410
Outlet_Location_Type  0
```

```

Outlet_Type          0
Item_Outlet_Sales    0
dtype: int64

```

data.dtypes

```

Item_Identifier      object
Item_Weight          float64
Item_Fat_Content      object
Item_Visibility      float64
Item_Type            object
Item_MRP             float64
Outlet_Identifier     object
Outlet_Establishment_Year  int64
Outlet_Size          object
Outlet_Location_Type  object
Outlet_Type          object
Item_Outlet_Sales    float64
dtype: object

```

data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       8523 non-null   object
1   Item_Weight                          7060 non-null   float64
2   Item_Fat_Content                     8523 non-null   object
3   Item_Visibility                      8523 non-null   float64
4   Item_Type                            8523 non-null   object
5   Item_MRP                            8523 non-null   float64
6   Outlet_Identifier                    8523 non-null   object
7   Outlet_Establishment_Year            8523 non-null   int64
8   Outlet_Size                          6113 non-null   object
9   Outlet_Location_Type                 8523 non-null   object
10  Outlet_Type                          8523 non-null   object
11  Item_Outlet_Sales                    8523 non-null   float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.2+ KB

```


data.describe()

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8523
mean	12.857645	0.066132	140.992782	1997.831867	2181
std	4.643456	0.051598	62.275067	8.371760	1706

check the correlation in columns

```
sns.heatmap(data.corr())
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fd946641ad0>



Box plot of understanding distribution and check for outliers

```
sns.boxplot(x="Item_Fat_Content", y="Item_Outlet_Sales", data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f835a82c250>
```



```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       8523 non-null   object
1   Item_Weight                           7060 non-null   float64
2   Item_Fat_Content                       8523 non-null   object
3   Item_Visibility                       8523 non-null   float64
4   Item_Type                             8523 non-null   object
5   Item_MRP                              8523 non-null   float64
6   Outlet_Identifier                     8523 non-null   object
7   Outlet_Establishment_Year             8523 non-null   int64
8   Outlet_Size                           6113 non-null   object
9   Outlet_Location_Type                  8523 non-null   object
10  Outlet_Type                           8523 non-null   object
11  Item_Outlet_Sales                     8523 non-null   float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.2+ KB
```

Replace the null values with median for numeric columns

```
data['Item_Weight'] = data['Item_Weight'].fillna(data['Item_Weight'].median())
```

```
data.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          0
Item_Fat_Content     0
Item_Visibility      0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size          2410
Outlet_Location_Type  0
Outlet_Type          0
Item_Outlet_Sales    0
dtype: int64
```

Replace the null values with most occuring data for non numeric columns

```
data['Outlet_Size'] = data['Outlet_Size'].fillna(data.groupby('Outlet_Size')['Item_Identifier
```

```
data.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          0
Item_Fat_Content     0
Item_Visibility      0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size          0
Outlet_Location_Type 0
Outlet_Type          0
Item_Outlet_Sales    0
dtype: int64
```

Encode all non numeric columns

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
columns = ['Outlet_Type', 'Outlet_Location_Type', 'Outlet_Size', 'Outlet_Identifier', 'Item_Type']
for column in columns:
    data[column] = encoder.fit_transform(data[column])
```

```
from sklearn.model_selection import train_test_split
```

```
x = data.drop(['Item_Outlet_Sales'], axis = 1)
y = data['Item_Outlet_Sales']
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

```
model = RandomForestRegressor()
```

```
model.fit(x,y)
```

```
RandomForestRegressor()
```


Readint the test data

```
test_data = pd.read_csv('/content/test_AbJTz21.csv')
```

```
test_data1 = pd.read_csv('/content/test_AbJTz21.csv')
```

```
test_data1.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	FDW58	20.750	Low Fat	0.007565	Snack Foods	107.8621
1	FDW14	8.300	reg	0.038428	Dairy	87.3198
2	NCN55	14.600	Low Fat	0.099575	Others	241.7538
3	FDQ58	7.315	Low Fat	0.015388	Snack Foods	155.0340
4	FDY38	NaN	Regular	0.118599	Dairy	234.2300



```
test_data.isnull().sum()
```

```
Item_Identifier      0
Item_Weight          0
Item_Fat_Content     0
Item_Visibility      0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size          0
Outlet_Location_Type 0
Outlet_Type          0
dtype: int64
```

Transforming the test data

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
columns = ['Outlet_Type', 'Outlet_Location_Type', 'Outlet_Size', 'Outlet_Identifier', 'Item_Type']
for column in columns:
    test_data[column] = encoder.fit_transform(test_data[column])
```

```
test_data['Outlet_Size'] = test_data['Outlet_Size'].fillna(test_data.groupby('Outlet_Size')['Item_Weight'].median())
test_data['Item_Weight'] = test_data['Item_Weight'].fillna(test_data['Item_Weight'].median())
```

```
test_data.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	1103	20.750	1	0.007565	13	107.8622
1	1067	8.300	4	0.038428	4	87.3198
2	1406	14.600	1	0.099575	11	241.7538
3	809	7.315	1	0.015388	13	155.0340
4	1184	12.500	2	0.118599	4	234.2300



```
y_pred = model.predict(test_data)
```

```
pred_data = pd.DataFrame(y_pred)
pred_data.to_csv('test.csv')
```

```
pred_data.head()
```

	Item_Identifier	Outlet_Identifier	Item_Outlet_Sales
0	FDW58	OUT049	1785.109670
1	FDW14	OUT017	1401.402472
2	NCN55	OUT010	693.736968
3	FDQ58	OUT017	2385.834378
4	FDY38	OUT027	6109.167744



Giving the proper column name for the predicted values

```
pred_data['Item_Identifier'] = test_data1['Item_Identifier']
```

```
pred_data['Outlet_Identifier'] = test_data1['Outlet_Identifier']
```

```
pred_data = pred_data.rename(columns={0: 'Item_Outlet_Sales'})
```

```
pred_data = pred_data[['Item_Identifier', 'Outlet_Identifier', 'Item_Outlet_Sales']]
```

find the current parameters

```
print('current parameters', model.get_params())
```

```
current parameters {'bootstrap': True, 'ccp_alpha': 0.0, 'criterion': 'squared_error',
```

```
from sklearn.model_selection import RandomizedSearchCV
```

creating random values for hyper parameters and creating grid

```
# Number of trees in random forest
n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
```

```
# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}
```

```
print(random_grid)
```

```
{'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000], 'max_feature
```

```
# Use the random grid to search for best hyperparameters
# First create the base model to tune
rf = RandomForestRegressor()
# Random search of parameters, using 3 fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_iter = 100)
# Fit the random search model
rf_random.fit(x, y)
```

```
Fitting 3 folds for each of 10 candidates, totalling 30 fits
/usr/local/lib/python3.7/dist-packages/joblib/externals/loky/process_executor.py:705: UserWarning:
  "timeout or by a memory leak.", UserWarning
RandomizedSearchCV(cv=3, estimator=RandomForestRegressor(), n_jobs=-1,
```



```

param_distributions={'bootstrap': [True, False],
                    'max_depth': [10, 20, 30, 40, 50, 60,
                                   70, 80, 90, 100, 110,
                                   None],
                    'max_features': ['auto', 'sqrt'],
                    'min_samples_leaf': [1, 2, 4],
                    'min_samples_split': [2, 5, 10],
                    'n_estimators': [200, 400, 600, 800,
                                     1000, 1200, 1400, 1600,
                                     1800, 2000]},

random_state=42, verbose=2)

```

print the best parameters

```
rf_random.best_params_
```

```

{'bootstrap': True,
 'max_depth': 10,
 'max_features': 'sqrt',
 'min_samples_leaf': 2,
 'min_samples_split': 5,
 'n_estimators': 200}

```

evaluating the best model

```

best_model = rf_random.best_estimator_
pred_tune = best_model.predict(test_data)

```

The best model is the model with default hyper parameter

```
pred_data_tuned = pred_data[['Item_Identifier', 'Outlet_Identifier', 'Item_Outlet_Sales']]
```

```
pred_data_tuned['Item_Outlet_Sales'] = pred_tune
```

```
pred_data_tuned.head()
```

```
pred_data_tuned.reset_index(drop=True, inplace=True)
pred_data.reset_index(drop=True, inplace=True)
```

	Item_Identifier	Outlet_Identifier	Item_Outlet_Sales
1	FDW14	OUT017	1394.587516

▼ predicted value in the test data for the normal model

```
pred_data.head()
```

	Item_Identifier	Outlet_Identifier	Item_Outlet_Sales
0	FDW58	OUT049	1785.109670
1	FDW14	OUT017	1401.402472
2	NCN55	OUT010	693.736968
3	FDQ58	OUT017	2385.834378
4	FDY38	OUT027	6109.167744

▼ predicted value in the test data for the tuned model.

the tuning was performed using RandomSearchCV

```
pred_data_tuned.head()
```

	Item_Identifier	Outlet_Identifier	Item_Outlet_Sales
0	FDW58	OUT049	1788.749769
1	FDW14	OUT017	1394.587516
2	NCN55	OUT010	817.322394
3	FDQ58	OUT017	2458.007444
4	FDY38	OUT027	5405.383431

lets see the score of each submission in Analytics Vidhya:

saving the both file into csv

```
pred_data.to_csv('test_normal.csv')
```

```
pred_data_tuned.to_csv('test_tuned.csv')
```

The score for Normal model is 1184.2392

The score for tuned model is 1163.8211

✓ 0s completed at 9:54 PM

